

# Intoduction\_to\_stocks\_r\_package

```
library(stocks_r)
```

## Overview

The **stocks\_r** package offers tools for preprocessing, analyzing, and modeling S&P 500 stock data organized by sector. It simplifies the end-to-end workflow of:

- **Loading and cleaning** historical S&P 500 price data
  - **Computing sector-level indices** from company-level stock prices
  - **Generating lag features and technical indicators** for forecasting
  - **Training Random Forest models** to predict sector index movements
  - **Evaluating model performance** using RMSE, MAE, and  $R^2$
- 

## Key Features

- **Sector-wise grouping**  
Automatically maps companies to sectors like Technology, Financials, or Health Care.
- **Technical indicator generation**  
Adds moving averages, returns, and momentum indicators for each sector's time series.
- **Random Forest pipeline**  
Trains models using default or customizable parameters via the **caret** package.
- **Performance comparison**  
Computes error metrics (RMSE, MAE,  $R^2$ ) to benchmark models across different sectors.
- **Visualization**  
Generates plots showing predicted vs actual sector index values and feature importance rankings.

## Overview

The `load_and_preprocess_data()` function in the **stocks\_r** package is designed to load historical S&P 500 stock price data from an Excel file and organize it into **sector-wise groups**.

Each sector is returned as a separate data frame, containing: - One column for **Date** - One column per company within that sector

This function is especially useful for: - Performing **sector-level financial analysis** - Building **machine learning models** on grouped company data

---

## What the Function Does

The function performs the following steps:

- **Reads** an Excel file where:
  - Rows represent daily price observations
  - Columns represent individual companies
- **Maps** each company to its corresponding sector (*e.g.*, *Energy*, *Health Care*)
- **Creates** a data frame for each sector that includes:
  - **Date** as the first column
  - Stock prices of the companies in that sector
- **Returns** a named list of sector-specific data frames

```
# Load and preprocess
sectors_data <- load_and_preprocess_data(df_cleaned2)

# View available sectors
names(sectors_data)
#> [1] "Energy"      "Health Care"
```

## View a Sample Sector

```
# Preview the first few rows of the Energy sector
head(sectors_data$Energy)
#>      Date chevron conocophillips schlumberger
#> 1 2010-01-01    76.99         38.9317        65.09
#> 2 2010-01-04    79.06         40.0828        67.11
#> 3 2010-01-05    79.62         40.1209        67.30
#> 4 2010-01-06    79.63         40.4106        68.80
#> 5 2010-01-07    79.33         40.2505        69.51
#> 6 2010-01-08    79.47         40.6012        70.65
```

## Sector Mapping Reference

The `load_and_preprocess_data()` function internally maps companies to the following sectors:

- **Energy:** Exxon Mobil, Chevron, ConocoPhillips, Schlumberger, EOG Resources, etc.
- **Health Care:** Johnson & Johnson, UnitedHealth Group, Eli Lilly, Pfizer, Merck & Company, etc.

## Return Value

The function returns a **named list of data frames**, where each name corresponds to a sector.

Each sector data frame contains:

- **Date:** The trading date
- Stock prices for all companies in that sector (one column per company)

## Inspect Sector Data Structure

```
# Check the structure of the Energy sector data
str(sectors_data$Energy)
#> 'data.frame':    2870 obs. of  4 variables:
#> $ Date           : Date, format: "2010-01-01" "2010-01-04" ...
#> $ chevron        : num  77 79.1 79.6 79.6 79.3 ...
#> $ conocophillips: num  38.9 40.1 40.1 40.4 40.3 ...
#> $ schlumberger   : num  65.1 67.1 67.3 68.8 69.5 ...
```

## Calculating Sector Index from S&P 500 Data

The `calculate_sector_index()` function computes the **average stock price across all companies in a sector** for each date, creating a simple sector-level time series. This is especially useful as input for forecasting or machine learning models.

### How It Works

- Takes a **sector-specific data frame** (e.g., `sectors_data$Energy`) created using `load_and_preprocess_data()`
- Handles missing values using linear interpolation (`na.approx`) and bidirectional filling (`fill`)
- Computes a row-wise average of all company columns
- Returns a data frame with:
  - `Date`: The trading date
  - `sector_index`: The average of all company prices in that sector

### Apply to this Data

```
# Compute sector index for the Energy sector
tech_index <- calculate_sector_index(sectors_data$Energy)

# View the result
head(tech_index)
#>      Date sector_index
#> 1 2010-01-01    60.33723
#> 2 2010-01-04    62.08427
#> 3 2010-01-05    62.34697
#> 4 2010-01-06    62.94687
#> 5 2010-01-07    63.03017
#> 6 2010-01-08    63.57373
```

## Calculating Technical Indicators for Sector Index

The `calculate_technical_indicators()` function in the `stocksr` package enriches a sector index time series with widely used technical indicators. These features are valuable for capturing patterns like momentum, trend, and volatility—key components for predictive modeling in finance.

### What Are Technical Indicators?

The function calculates the following indicators:

**\*Daily Returns\*\*** The daily return  $R_t$  is computed as the percent change from the previous day's index:

$$R_t = \frac{P_t - P_{t-1}}{P_{t-1}}$$

Where: -  $P_t$  is the sector index at time  $t$  -  $R_t$  is the return on day  $t$

**\*Simple Moving Averages (SMA)\*\*** The SMA over  $n$  days is:

$$\text{SMA}_t = \frac{1}{n} \sum_{i=0}^{n-1} P_{t-i}$$

Used for smoothing trends over time. The function includes: - SMA over 5, 10, and 20 days

**\*Exponential Moving Averages (EMA)\*\*** The EMA assigns more weight to recent prices. It's calculated recursively as:

$$\text{EMA}_t = \alpha \cdot P_t + (1 - \alpha) \cdot \text{EMA}_{t-1}$$

Where  $\alpha = \frac{2}{n+1}$  is the smoothing factor.

The function includes: - EMA over 5 and 10 days

**Rolling Volatility** The volatility (standard deviation) over a rolling window of  $n$  days:

$$\text{Volatility}_t = \sqrt{\frac{1}{n-1} \sum_{i=0}^{n-1} (P_{t-i} - \bar{P})^2}$$

Computed for 5 and 10-day windows.

**Relative Strength Index (RSI)** The RSI measures the magnitude of recent gains and losses:

$$\text{RSI} = 100 - \left( \frac{100}{1 + \frac{\text{Average Gain}}{\text{Average Loss}}} \right)$$

Used to identify overbought ( $\text{RSI} > 70$ ) or oversold ( $\text{RSI} < 30$ ) conditions.

**Lagged Features** For time series models, the function creates lagged values:

- $\text{Lag}_k = P_{t-k}$
- $\text{Returns\_Lag}_k = R_{t-k}$  for  $k = 1$  to 5

These help capture short-term dependencies and trends.

**Apply to Sector Data**

```
tech_features <- calculate_technical_indicators(tech_index)
head(tech_features)
```

#>	Date	sector_index	log_price	returns	log_returns	MA5	EMA5
#> 1	2010-03-11	58.99900	4.077520	-0.002252580	-0.002255121	59.05925	58.87591
#> 2	2010-03-12	59.21890	4.081241	0.003727182	0.003720253	59.13379	58.99024
#> 3	2010-03-15	59.08460	4.078970	-0.002267857	-0.002270433	59.09787	59.02170
#> 4	2010-03-16	59.75673	4.090282	0.011375779	0.011311561	59.23829	59.26671
#> 5	2010-03-17	60.55257	4.103512	0.013317886	0.013229982	59.52236	59.69533
#> 6	2010-03-18	60.06397	4.095410	-0.008069022	-0.008101753	59.73535	59.81821

#>	MA10	EMA10	MA20	EMA20	MA50	EMA50	MACD	MACD_signal
#> 1	58.29988	58.51287	57.98634	58.36753	59.52566	59.52566	-0.26033348	-0.9402428
#> 2	58.55539	58.64124	58.07375	58.44861	59.50329	59.51363	-0.13249964	-0.7786942
#> 3	58.74668	58.72185	58.15445	58.50918	59.44330	59.49680	-0.04924725	-0.6328048
#> 4	58.92152	58.91001	58.19737	58.62800	59.39149	59.50700	0.10773850	-0.4846962
#> 5	59.17667	59.20866	58.29412	58.81129	59.34361	59.54800	0.33676256	-0.3204044
#> 6	59.39730	59.36417	58.35258	58.93059	59.28428	59.56823	0.44567959	-0.1671876

#>	volatility5	volatility20	lag_1	log_return_lag_1	lag_2	log_return_lag_2
#> 1	0.008368733	0.010355585	59.13220	0.001312047	59.05467	-0.003541283
#> 2	0.004339017	0.010162117	58.99900	-0.002255121	59.13220	0.001312047
#> 3	0.003020810	0.010192372	59.21890	0.003720253	58.99900	-0.002255121
#> 4	0.005608456	0.008966452	59.08460	-0.002270433	59.21890	0.003720253
#> 5	0.007320907	0.009283416	59.75673	0.011311561	59.08460	-0.002270433
#> 6	0.008994578	0.009498558	60.55257	0.013229982	59.75673	0.011311561

#>	lag_5	log_return_lag_5	day_of_week	month
#> 1	57.85763	-0.002477155	Thu	Mar
#> 2	58.84620	0.016941865	Fri	Mar
#> 3	59.26417	0.007077591	Mon	Mar
#> 4	59.05467	-0.003541283	Tue	Mar
#> 5	59.13220	0.001312047	Wed	Mar
#> 6	58.99900	-0.002255121	Thu	Mar

These engineered features are ready to be used in time series forecasting, machine learning pipelines, or financial dashboarding.

## Random Forest Modeling for Sector Index Prediction

The `build_random_forest_model()` function trains a **Random Forest regression model** to predict a sector's daily index using engineered features such as returns, lag values, and technical indicators. It evaluates the model using standard performance metrics and provides feature importance to interpret the model.

### What is a Random Forest?

A **Random Forest** is an ensemble machine learning algorithm that builds multiple decision trees and combines their predictions. In regression tasks, it averages the results from each tree to provide a robust, stable prediction.

### Why use Random Forests for financial data?

- Handles non-linear relationships between features and response
- Robust to noise and overfitting

- Provides built-in feature importance for interpretation
- Performs well on datasets with many engineered variables (like technical indicators)

## How the Function Works

The `build_random_forest_model()` function performs the following steps:

1. **Input:** A data frame with:
  - Date
  - `sector_index` (target)
  - Engineered features (technical indicators, returns, lags, etc.)
2. **Split the data:**
  - 80% for training
  - 20% for testing
  - Chronological order is preserved to respect time-dependence
3. **Train the model:**
  - A Random Forest is trained on the training set using `randomForest::randomForest()`
  - Parameters include `ntree = 100`, `mtry = sqrt(p)`, `nodesize = 5`
4. **Evaluate:**
  - Calculates:
    - **RMSE** (Root Mean Squared Error)
    - **MAE** (Mean Absolute Error)
    - $R^2$  (Coefficient of Determination)
  - Plots actual vs predicted values
5. **Feature Importance:**
  - Displays top variables ranked by `IncNodePurity`

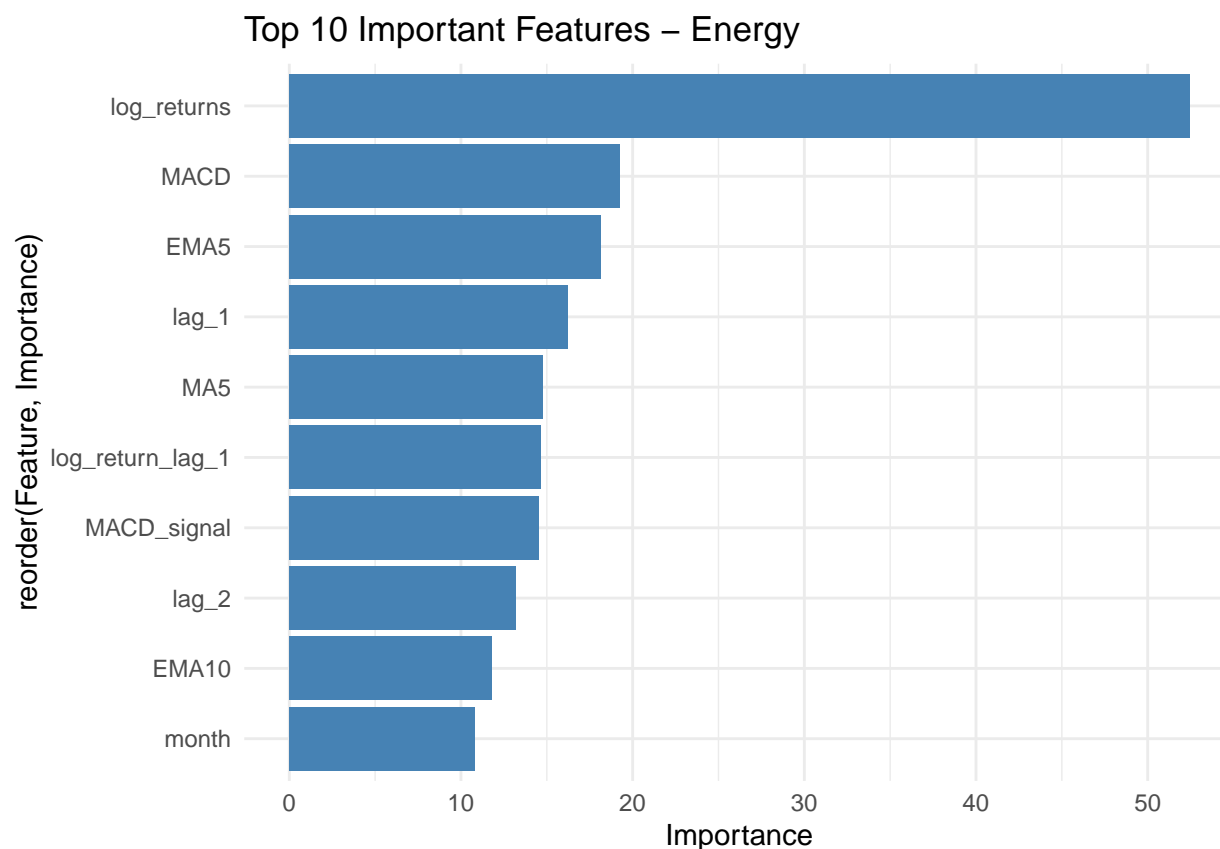
## Apply to Your Data

```
# Train and evaluate the model
rf_results <- build_random_forest_model(tech_features, "Energy")
#>
#> --- Building Random Forest model for Energy sector ---
#> Model performance:
#> RMSE: 0.7691
#> MAE: 0.5274
#> R<U+00B2>: 0.9963
```

## Energy Sector – Random Forest Predictions

$R^2 = 0.996$





## Output

The function returns a named list:

Name	Description
model	The trained <b>randomForest</b> object
rmse	Root Mean Squared Error on the test set
mae	Mean Absolute Error on the test set
r2	R-squared score on the test set
feature_importance	Data frame ranking features by <b>IncNodePurity</b>

This modeling function is a key part of the **stocksr** pipeline and can be extended to all sectors for comparative performance analysis.

```
# Step 3: View results
rf_results$rmse
#> [1] 0.7690529
rf_results$r2
#> [1] 0.9963346
head(rf_results$feature_importance)
#> NULL
```